

Kentaro Vadney

Alameda, CA 94502

(424) 533-9384

kentarovadney@gmail.com

linkedin.com/in/kentarov

December 11, 2024

Hiring Manager

Apple Inc.

Remote Position - Python Developer (ML Pipelines)

Dear Hiring Manager,

I am writing to express my strong interest in the Python Developer – ML Pipelines position at Apple. With over three years of expert-level Python development experience and a proven track record building scalable ML infrastructure, I am excited about the opportunity to contribute to Apple's world-class machine learning systems.

Technical Excellence in ML Pipeline Development

Throughout my career, I have specialized in designing and optimizing production ML pipelines that handle large-scale data processing. At Centigrade, I architected high-performance API services using Python and FastAPI that supported a 5x increase in query throughput while reducing database latency by 60%. I built an AWS Kendra AI search system with a Python-based indexing pipeline processing over 10,000 documents, demonstrating my ability to create robust, scalable ML workflows.

Currently, as a Machine Learning Engineer at Mercor (\$100/hour contract role), I design and implement distributed ML pipelines for candidate matching systems using PyTorch and advanced distributed computing frameworks. This experience has given me deep expertise in building production-grade data processing workflows that maintain reliability at scale.

Distributed Computing & System Optimization

My experience with distributed systems aligns perfectly with Apple's need for Ray expertise and scalable computing. As a researcher at UC Berkeley's RISE Lab, I led a team investigating consensus protocols and distributed system reliability, co-authoring peer-reviewed research on Byzantine fault tolerance optimizations. At Belden, I developed ML-powered threat detection pipelines deployed on Kubernetes, reducing inference latency by 40% through algorithmic refinement and parallel processing.

I have extensive hands-on experience with:

- **Distributed frameworks:** Dask, Kubernetes orchestration
- **Performance optimization:** Algorithm design, parallel processing, caching strategies
- **Cloud platforms:** AWS (EC2, S3, Lambda), GCP, Azure with Terraform infrastructure-as-code
- **Production systems:** CI/CD pipelines, automated testing (pytest), version control (Git)

Collaborative Problem-Solving

I excel at translating complex ML research into production systems through collaboration with data scientists and engineers. At CarbonSustain, I built distributed transaction processing systems with Python microservices handling 1,000+ daily transactions across three enterprise deployments. At Belden, I enhanced threat detection models by 30% through iterative feature engineering and close collaboration with security researchers.

Why Apple

Apple's commitment to pushing the boundaries of machine learning at scale resonates deeply with my passion for building robust, efficient ML infrastructure. I am particularly drawn to the opportunity to work on large-scale ML pipelines that power products used by millions of people worldwide. The challenge of optimizing compute efficiency and system reliability at Apple's scale is exactly the type of complex, impactful problem I thrive on solving.

I am available to start within two weeks' notice and would welcome the opportunity to discuss how my expertise in Python development, distributed computing, and ML pipeline optimization can contribute to Apple's mission. Thank you for considering my application.

Sincerely,

Kentaro Vadney