STATEMENT OF QUALIFICATIONS

Kentaro Vadney
Database Administrator Position

Describe your experience with Oracle database installation, configuration, and upgrade. Include specific versions and environments you have worked with.

While I do not have direct experience with Oracle database systems, I bring over two years of hands-on production database engineering experience with enterprise-grade relational databases, specifically PostgreSQL and MySQL. My work has encompassed the complete database lifecycle including installation, configuration, schema design, performance optimization, and maintenance in production environments serving enterprise clients.

At Centigrade, I led PostgreSQL optimization initiatives that required deep understanding of database architecture, configuration tuning, and schema redesign. I reduced query latency by 60% and enabled 5x throughput increases through systematic performance analysis, index optimization, and configuration adjustments. This work required intimate knowledge of database internals, query execution planning, and system resource management—skills that translate directly to Oracle database administration.

During my tenure at CarbonSustain as Founding Engineer, I architected production PostgreSQL databases from the ground up for three enterprise pilot programs. This involved selecting appropriate versions, configuring instances for optimal performance based on workload characteristics, designing backup strategies, and establishing monitoring frameworks. I worked extensively with AWS RDS PostgreSQL in versions 13 through 15, managing database configurations, parameter tuning, and version compatibility considerations during platform evolution.

I have also maintained MySQL databases in production environments and worked with MongoDB for NoSQL use cases, demonstrating my ability to quickly master different database technologies and their unique architectural considerations. My experience spans Linux-based database deployments, containerized database instances using Docker and Kubernetes, and cloud-managed database services on AWS.

I am committed to rapidly acquiring Oracle-specific expertise and am confident that my strong foundation in relational database systems, combined with my proven track record of mastering complex technical systems, positions me to quickly become proficient with Oracle database administration.

Describe your experience with Oracle RAC, Data Guard, or other high-availability/disaster recovery solutions. What were the challenges and how did you address them?

While I have not worked directly with Oracle RAC or Data Guard, I have extensive experience designing and implementing high-availability and disaster recovery solutions for production database systems using analogous technologies and architectural patterns.

At CarbonSustain, I designed comprehensive backup and recovery solutions for mission-critical PostgreSQL databases supporting enterprise clients. I implemented automated snapshot systems with point-in-time recovery capabilities, achieving a 2-hour Recovery Time Objective (RTO) and 15-minute Recovery Point Objective (RPO). This required careful planning around backup windows, storage capacity management, transaction log archiving, and testing recovery procedures under various failure scenarios.

I configured high-availability PostgreSQL deployments utilizing streaming replication and failover mechanisms to ensure service continuity. This work involved understanding replication lag, synchronous versus asynchronous replication trade-offs, and automated failover procedures—concepts that directly parallel Oracle Data Guard's standby database architecture.

The primary challenges I encountered included:

Challenge 1: Minimizing replication lag while maintaining primary database performance. I addressed this by carefully tuning replication parameters, optimizing network bandwidth between primary and standby systems, and implementing monitoring to alert on replication delays exceeding acceptable thresholds.

Challenge 2: Ensuring backup integrity and testing recovery procedures without impacting production. I established automated backup validation processes and created isolated environments for conducting regular disaster recovery drills, documenting procedures and improving recovery time with each iteration.

Challenge 3: Balancing high availability requirements with cost constraints. I designed tiered backup strategies distinguishing between hot standbys for immediate failover and cold backups for long-term retention, optimizing both recovery capabilities and infrastructure costs.

At Centigrade, I maintained critical backend infrastructure with 99.9% uptime SLA requirements, providing 24/7 production support and incident response. I troubleshot complex database performance issues, connection pooling problems, and system degradation scenarios, often under time pressure with direct client impact. This operational experience has prepared me for the demands of maintaining enterprise-grade high-availability database systems.

My background in distributed systems, cloud infrastructure management with Terraform, and container orchestration with Kubernetes provides additional foundation for understanding and managing complex, redundant database architectures typical of Oracle RAC and Data Guard environments.

Describe your experience working in a team environment. How do you collaborate with developers, system administrators, and business analysts?

Throughout my career, I have served as a technical bridge between development teams, infrastructure operations, and business stakeholders, consistently fostering collaborative relationships that drive project success.

At Centigrade, I served as technical lead maintaining critical backend infrastructure and database systems, which required constant collaboration across multiple teams. I worked directly with software developers to optimize their database queries, redesign schemas for improved performance, and establish best practices for database interactions. When developers encountered performance issues, I performed query analysis, explained execution plans, and recommended indexing strategies or query rewrites. I made complex database concepts accessible to developers with varying levels of database expertise, ensuring they understood not just what to change, but why those changes would improve system performance.

My collaboration with system administrators involved coordinating database deployments, troubleshooting connectivity issues, managing backup storage, and ensuring proper monitoring integration. I established CI/CD pipelines using GitHub Actions and CircleCI that required close coordination with DevOps teams to ensure database schema migrations were safely automated alongside application deployments. When incidents occurred, I worked alongside system administrators to diagnose whether issues stemmed from database configuration, infrastructure resources, or application behavior.

As Founding Engineer at CarbonSustain, I led the engineering team and served as primary technical liaison with business stakeholders. I translated business requirements from pilot clients into technical database architectures, explaining trade-offs between different approaches in terms business analysts could understand—focusing on reliability, performance, and cost implications rather than technical implementation details. I regularly provided status updates on infrastructure projects, database performance metrics, and capacity planning to non-technical stakeholders.

My approach to collaboration emphasizes:

Clear Communication: I document technical decisions, write comprehensive runbooks for common procedures, and maintain updated system diagrams. During my time at CarbonSustain, I created documented resolution procedures for critical incidents, enabling team members to respond effectively during off-hours.

Proactive Problem-Solving: Rather than waiting for issues to escalate, I implemented comprehensive monitoring systems tracking database performance metrics, connection pools, and resource utilization. This allowed me to identify and address potential problems before they impacted users or other teams.

Knowledge Sharing: I believe in elevating the entire team's technical capabilities. I've conducted informal training sessions on database optimization techniques, created documentation for common troubleshooting procedures, and made myself available for technical questions from team members.

Flexibility and Responsiveness: In supporting production systems with 99.9% uptime SLAs and providing 24/7 support, I've demonstrated reliability and responsiveness that builds trust with colleagues. I understand that database issues can block multiple teams, and I prioritize rapid communication and resolution.

My experience leading technical teams, managing cross-functional projects, and serving as the database subject matter expert has prepared me to collaborate effectively with diverse stakeholders in a team-oriented database administration role.